
CSE 546 Assignment 2 - Deep Q-Networks

Ijaj Ahmed

Department of Computer Science
University at Buffalo, SUNY
Buffalo, NY 14260
ijajahme@buffalo.edu

Mohammed Nasheed Yasin

Department of Linguistics
University at Buffalo, SUNY
Buffalo, NY 14260
m44@buffalo.edu

Abstract

This report presents our experiments on a self-created GridWorld environment and two additional environments, CartPole-v1 and LunarLander-v2. These environments were selected from the ClassicalControl and Box2D collections in the Gymnasium library [2]. We applied DQN and Double DQN to solve these environments and conducted a case study on the outcomes.

Environment Description

1 The Grid-world Environment

There are certain features that are common to both the stochastic and deterministic environments.

Environmental Elements



Figure 1: Environmental elements (from left to right; top to bottom) Agent, Goal, Negative Reward, Reward

The environment is a 6×6 grid defined according to the Gym [2] API. With 36 possible positions that the agent can occupy. The goal is to reach the oasis (shown in Figure 1) **after consuming all** the juice (positive reward) within a (configurable) maximum number of time steps. If the agent lands on a juice tile it is awarded +0.99 and the cactus (negative reward) leads to a -1.0 reward. Once all the juice is consumed, the agent must proceed to the oasis to earn a reward of +1.0. The states in this environment are a **combination** of the agent's position and the currently available rewards (positive, negative and goal) on the grid. The Formula 1 gives us the number of possible states for the agent.

$$num_{states} = num_{pos} \sum_{k=0}^{c_{reward}} \binom{c_{reward}}{k} \quad (1)$$

Here num_{pos} refers to the number of grid squares, 36 in our case. c_{reward} refers to the number of positive rewards + the number of negative rewards + 1 (for the goal state). We have 6 negative rewards, 3 positive rewards and one goal. Hence, the num_{states} for us is 36864.

In each position our agent can take 4 potential actions: 1. Left 2. Right 3. Up 4. Down. Resetting the environment will not change the location or distribution of the rewards and goal state. It only alters the initial state of the agent.

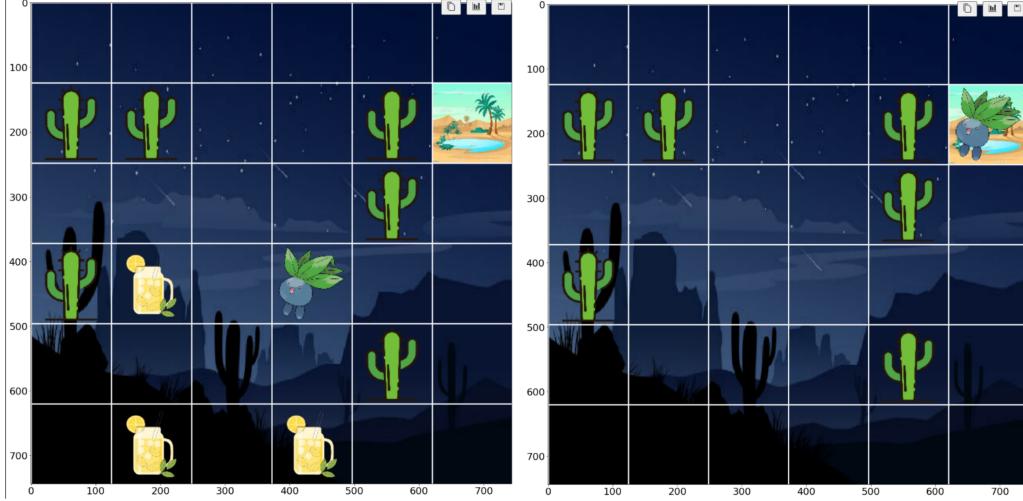


Figure 2: Environment visualizations

1.1 Safety in AI

The following are properties of the environment that ensure valid behavior from the agent:

1. We ensure that the agent consumes *all* the pos_reward and reaches the goal_state in the fewest number of steps by imposing a penalty of $0.1 \times$ the max_reward for every move made.
2. The reward on all squares is consumed once the agent lands in that state, preventing the agent from settling down in a *high-reward* neighborhood.
3. The result of any action (left, right, top, down) are clipped to the min and max values of 0 and GridSize (6 in our case) respectively, ensuring that the agent never leaves the environment.
4. If the agent makes a move but remains in the same spot, we impose the *maximum negative reward* (-1.0 in our case). This allows the agent to disincentivize making fruitless moves.
5. We also prevent a *goal rush* (before the agent consumes all the positive rewards) by making the goal square unreachable when there are positive reward squares left. If the agent takes

an action to move into a goal square before collecting all the *positive rewards*, they will be kept on the same square, incurring an additional penalty (-1.0 in our case) as detailed in the previous point.

6. The stochasticity of the starting point and limited time steps will nudge the agent to build strategies that accumulate the maximal reward in the shortest time.

2 CartPole-v1

In the cart-pole problem version described by Barto, Sutton, and Anderson [1] a pole is connected to a cart through a joint that cannot be moved. The cart can move on a track without any friction. The pole is positioned upright on the cart and the objective is to keep the pole balanced by applying forces to the cart in either the left or right direction. This environment has been visualized in Figure 3.

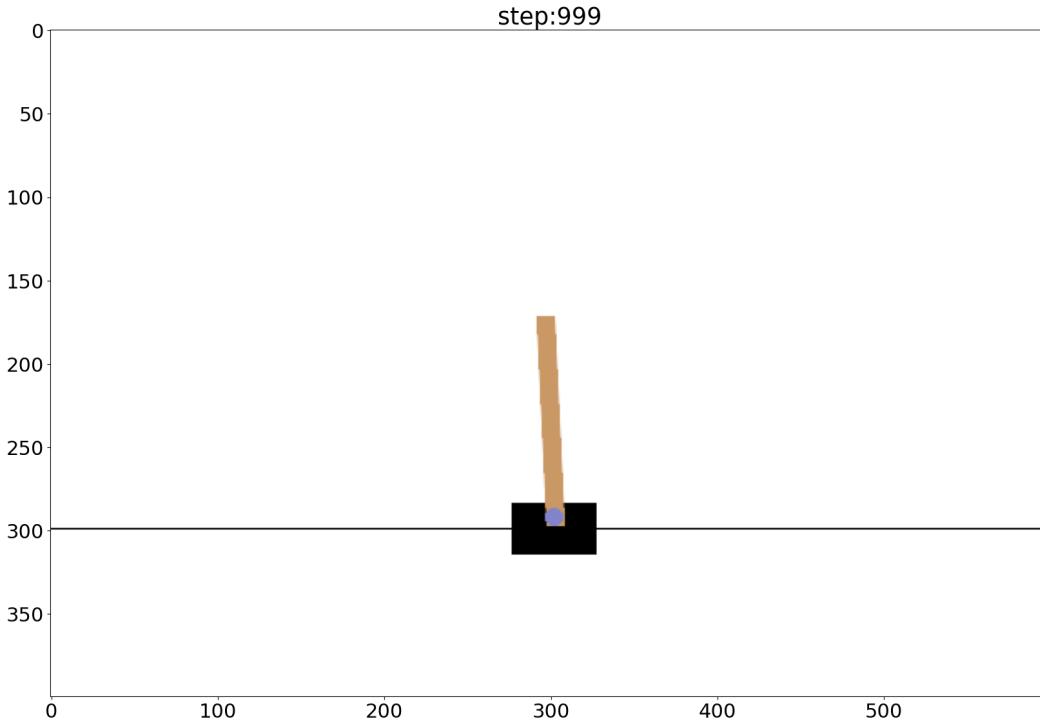


Figure 3: A frame from CartPole-v1

2.1 Action Space

The action is an ndarray of shape (1,) that can take on values 0, 1 to indicate the direction in which the cart is pushed with a fixed force.

- 0: The cart is pushed to the left
- 1: The cart is pushed to the right

The velocity that is either decreased or increased by the force applied is not constant and depends on the angle at which the pole is pointing. The center of gravity of the pole affects the amount of energy required to move the cart beneath it.

2.2 Observation Space

The observation is an ndarray of shape (4,) where the values represent the positions and velocities described next.

| Num | Observation | Min | Max |
|-----|-----------------------|--|--------------------------------------|
| 0 | Cart Position | -4.8 | 4.8 |
| 1 | Cart Velocity | -Inf | Inf |
| 2 | Pole Angle | $\sim -0.418 \text{ rad } (-24^\circ)$ | $\sim 0.418 \text{ rad } (24^\circ)$ |
| 3 | Pole Angular Velocity | -Inf | Inf |

Although the ranges above indicate the possible values for each element in the observation space, they do not reflect the allowed values of the state space in an ongoing episode. Specifically:

- The cart x-position (index 0) can be take values between $(-4.8, 4.8)$, but the episode terminates if the cart leaves the $(-2.4, 2.4)$ range.
- The pole angle can be observed between $(-.418, .418)$ radians or $(\pm 24^\circ)$, but the episode terminates if the pole angle is not in the range $(-.2095, .2095)$ or $(\pm 12^\circ)$

2.3 Rewards

The objective of the task is to keep the pole upright for as long as possible. To encourage this behavior, a reward of +1 is given for every step taken, including the final step when the episode terminates. In version 1 of the task, the threshold for achieving a successful outcome is set at 475.

2.4 Start State

All observations are assigned a uniformly random value in $(-0.05, 0.05)$

2.5 Episode End

The episode terminates under these conditions:

1. Termination: Pole Angle is greater than $\pm 12^\circ$
2. Termination: Cart Position is greater than ± 2.4 (center of the cart reaches the edge of the display)
3. Truncation: Episode length is greater than 500

3 LunarLander-v2

This environment represents a classic problem of optimizing rocket trajectory. Based on Pontryagin's maximum principle, the optimal approach is to either fire the engine at full throttle or turn it off completely. As a result, this environment has discrete actions: the engine is either on or off.

Two versions of the environment are available: discrete and continuous. In our work we have used the discrete version. The landing pad is always located at coordinates $(0, 0)$, which are represented by the first two numbers in the state vector. It is possible to land outside of the landing pad. Since fuel is unlimited, an agent can learn to fly and land successfully on its first attempt.

3.1 Observation Space

The state of the environment is represented by an 8-dimensional vector that includes the x and y coordinates of the lander, its linear velocities in x and y, its angle and angular velocity, and two boolean values indicating whether each leg is in contact with the ground.

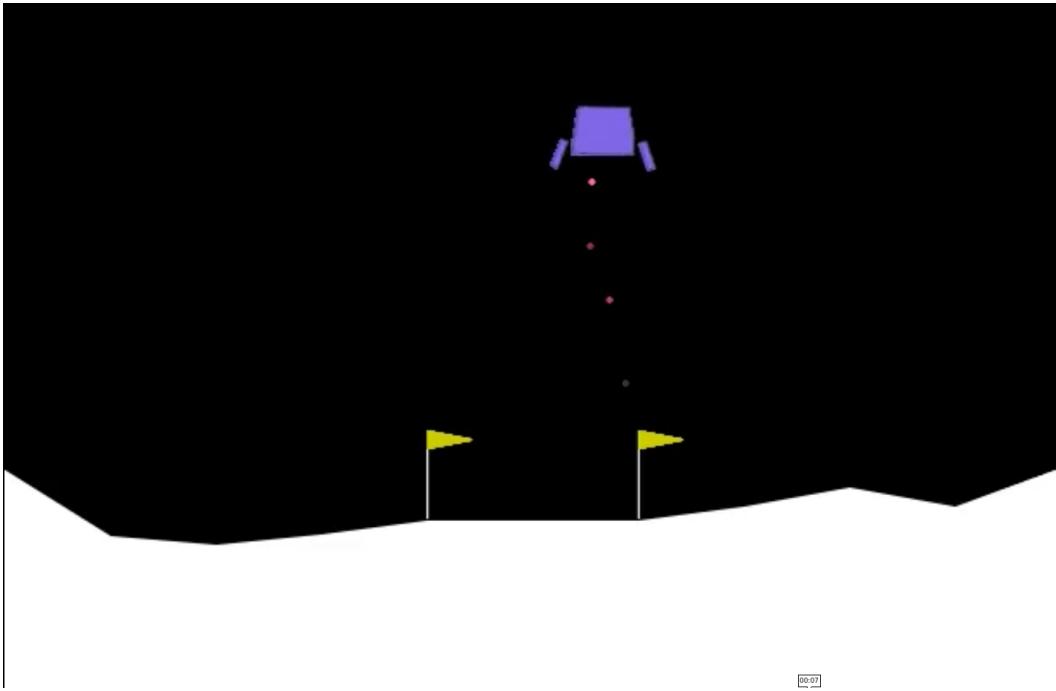


Figure 4: A frame from LunarLander-v2

3.2 Action Space

Four discrete actions are available in this environment: remain idle, activate the left orientation engine, activate the main engine, or activate the right orientation engine.

3.3 Reward

A reward is given after each step in the environment. The total reward for an episode is calculated by summing the rewards for all steps within that episode. The reward for each step is determined by the following factors:

- The reward increases/decreases as the lander gets closer/further from the landing pad.
- The reward increases/decreases as the lander moves slower/faster.
- The reward decreases as the lander tilts more (angle not horizontal).
- The reward increases by 10 points for each leg in contact with the ground.
- The reward decreases by 0.03 points for each frame a side engine is firing.
- The reward decreases by 0.3 points for each frame the main engine is firing.

An additional reward of -100 or +100 points is given for crashing or landing safely, respectively. An episode is considered solved if it scores at least 200 points.

3.4 Starting State

At the beginning of each episode, the lander is positioned at the top center of the viewport and a random initial force is applied to its center of mass.

3.5 Episode End

An episode terminates if any of the following conditions are met:

1. The lander crashes (its body comes into contact with the moon).
2. The lander moves outside of the viewport (its x coordinate is greater than 1).
3. The lander is not awake. According to the Box2D documentation, a body that is not awake does not move or collide with any other body.

4 Discuss the benefits of:

4.1 Using experience replay in DQN and how its size can influence the results

The experience replay in DQN helps to curb the problem of correlation between consecutive input samples to the training network which otherwise causes the network to diverge or make it unstable. The experience relay ensures the samples are i.i.d (independent and identically distributed) by randomizing them among batches by same distribution and are kept independent of each other in the same batch. This avoids the overfitting of some group of consecutive samples as the temporal variation is poor. The experience relay provides the opportunity to train the network by considering a large pool of samples already trained. There is a tradeoff between choosing the buffer size and the final outcome of such choice. Generally a small size buffer helps to learn the environment rapidly at the cost of ignoring the previous experiences. Similarly a larger buffer suffers from slow learning

4.2 Introducing the target network:

The target network helps to avoid oscillation in policy as conventional Q value update changes the current estimate and target value. If the target Q value is impacted by the current update of the Q value, the training chases a non stationary target and does not help much in learning efficiently. WIth the introduction of a target network which is updated with the primary network only after some particular interval helps to train the model effectively.

4.3 Representing the Q function as $q'(s, w)$

This solves the problem of dealing with larger state-action domains or complex networks as the neural network representation can now serve as function approximation. The design matrix now reduces to updating the weight of the network to minimize the loss of the estimate and target network.

5 Show and discuss your results after applying your DQN implementation on the three environments.

5.1 GridWorld

We represent the current state of the environment as a 36 dimensional vector of rank 1. Here each component of the vector has one of 3 values which indicate:

- 0.5: Agent location
- 1.0: Negative reward
- 0.99: Positive reward
- 1.0: Goal square

We feed this vector to a Feed-forward Network that will give us the estimated action-values. Then We take the argmax over these action-values. We consider the environments as solved when we have achieved a 100 episode running average of 2.1.

5.2 CartPole-v1 & LunarLander-v2

For both these environments we have rather similar action spaces (Discrete), hence the Feed-forward Network architechture used is also similar. The input of the network has the shape 1 x state-space output of network will 1 x action-space.

The environments are considered solved when:

CartPole-v1 Running average reward of 500 (100 episodes)

LunarLander-v2 Running average reward > 200 (100 episodes)

5.3 Training Discussion

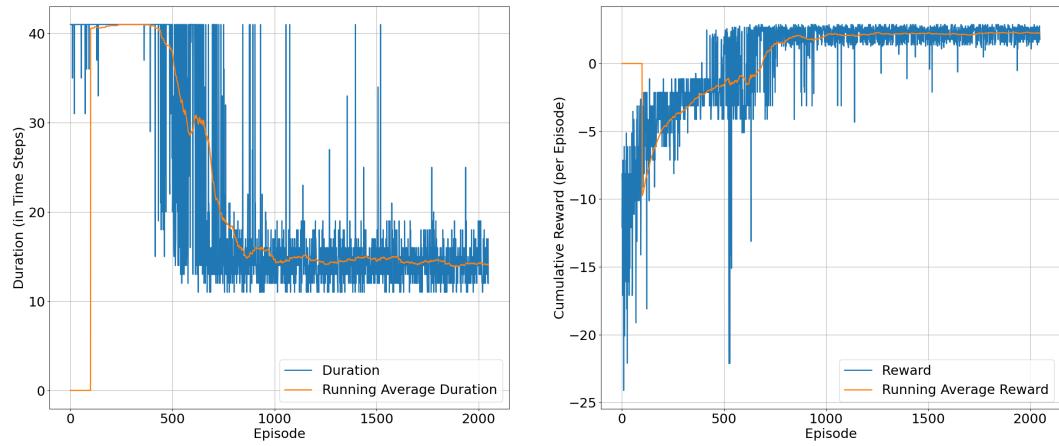


Figure 5: Grid Training

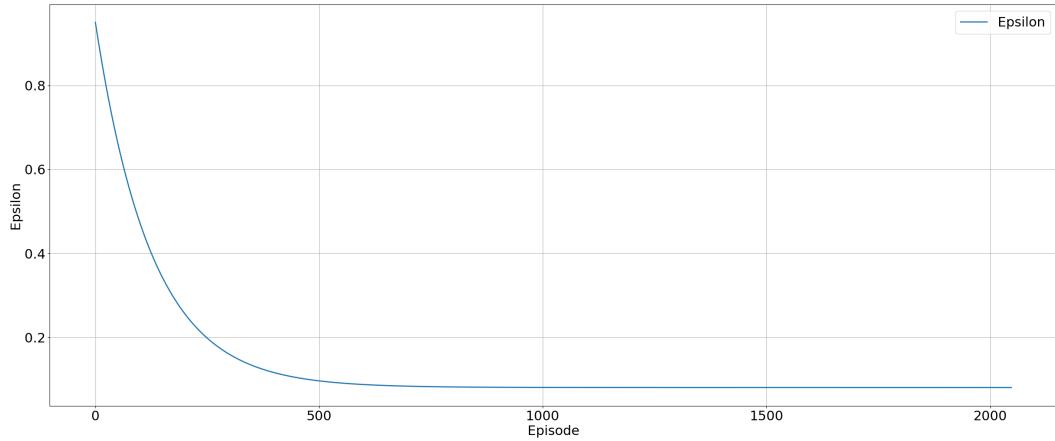


Figure 6: Epsilon Decay – Grid Env

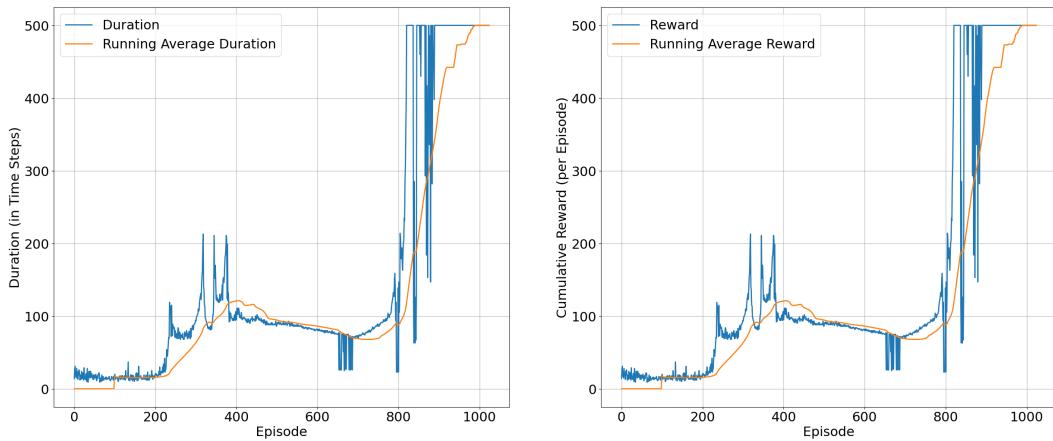


Figure 7: Cart-Pole Training

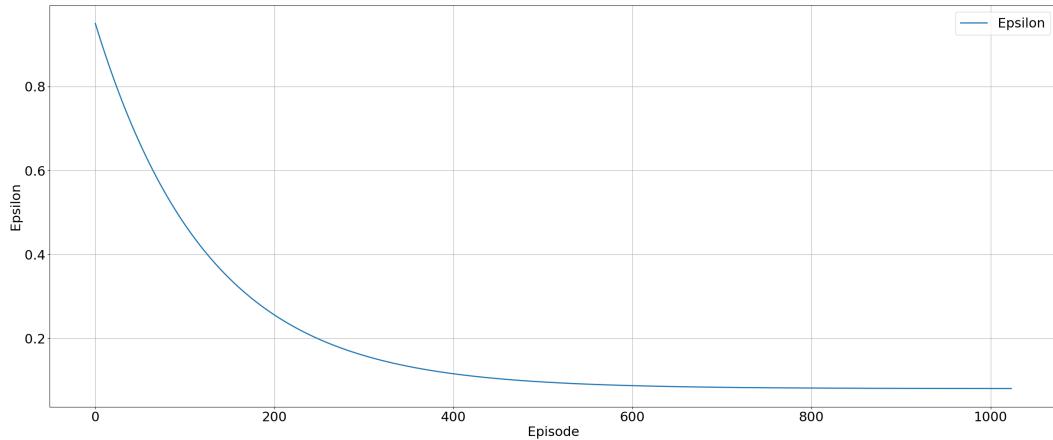


Figure 8: Epsilon Decay – Cart Pole

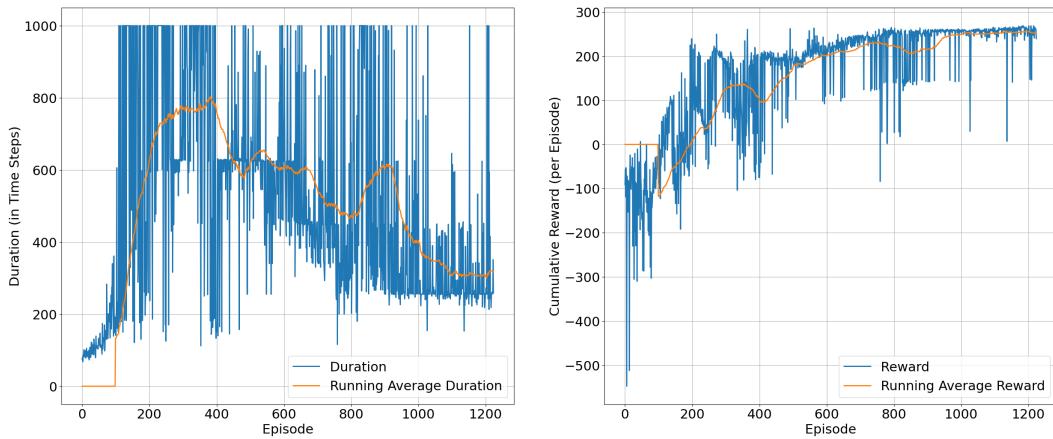


Figure 9: Lunar Lander Training

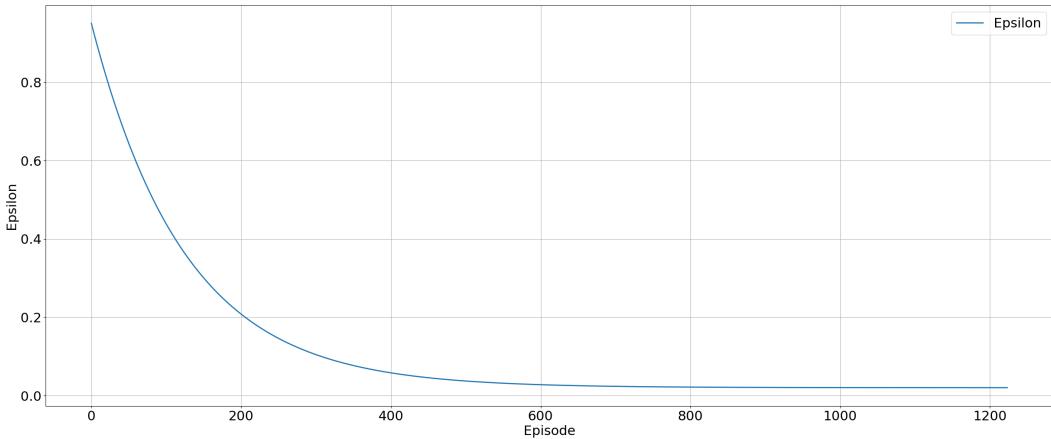


Figure 10: Epsilon Decay – Lunar Lander

We successfully solved all three environments using simple FeedForward neural networks. Additionally, we implemented a tau-based Target Network update strategy to stabilize the training process and used an experience replay memory of size 10000. Of the three environments, CartPole-v1 exhibited the most unstable learning. This instability is due to the fact that a single misstep in CartPole-v1 can cause the pole to collapse and the episode to end prematurely. The other two environments had more stable learning curves and were successfully solved by the DQN agent.

6 Provide the evaluation results. Run your agent on the three environments for at least 5 episodes, where the agent chooses only greedy actions from the learnt policy.

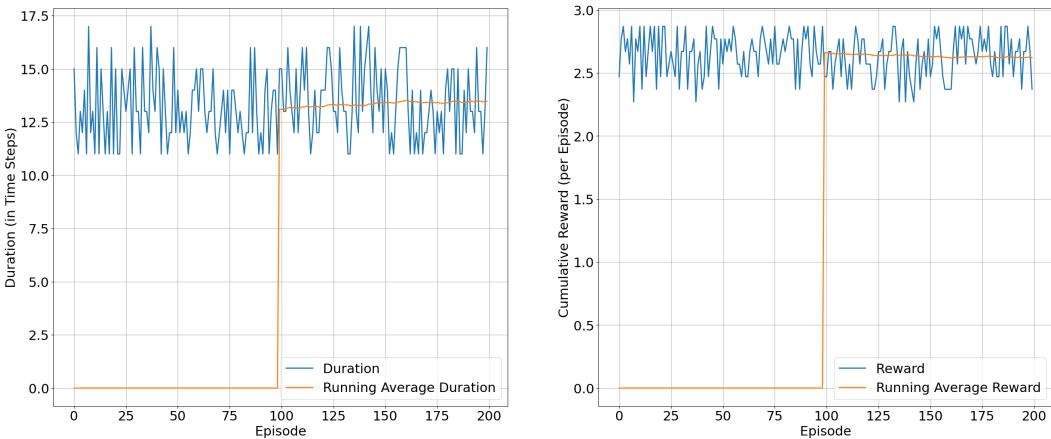


Figure 11: Grid Environment Evaluation

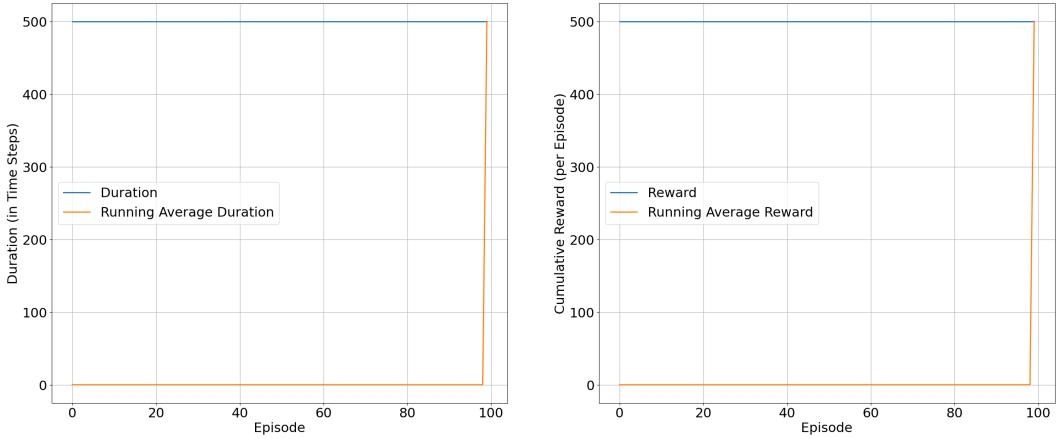


Figure 12: Cart Environment Evaluation

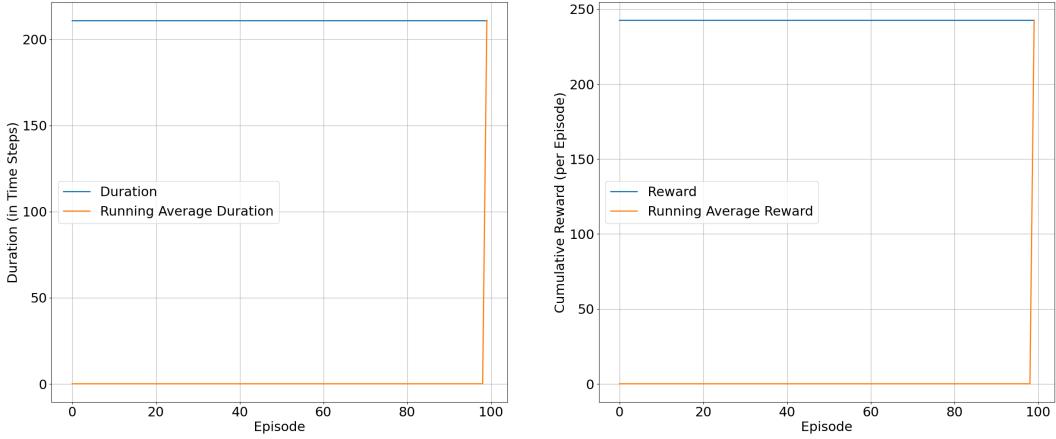


Figure 13: Lunar Environment Evaluation

The Deep Q-Learning [3] with a FeedForward neural network approach has successfully learned the optimal policy for all three environments. The evaluation result for the GridEnvironment varies due to the agent being placed at a random start squares, which affect the number of steps required to reach the goal. To encourage the agent to reach the goal as quickly as possible, we applied a negative reward of -0.1 for each step taken.

7 Discuss the Double Deep Q-Learning algorithm and the main improvements over vanilla DQN.

Double Deep Q-Network (DDQN) [4] is an extension of the Deep Q-Network (DQN) algorithm, which is a deep reinforcement learning technique used to learn a policy for an agent in an environment.

In DDQN, the idea is to address the issue of overestimation of Q-values, which can occur in the original DQN algorithm. This overestimation can lead to suboptimal or even unstable policies being learned.

To address this, DDQN uses a second network (known as the "target" network) to help estimate the Q-values used to update the primary network. The target network is a copy of the primary network, but with frozen parameters that are only periodically updated. This helps to prevent the

overestimation of Q-values, as the target network provides a more accurate estimate of the true Q-values.

8 Show and discuss your results after applying your DDQN implementation on the environment. Plots should include epsilon decay and the total reward per episode

8.1 GridWorld

We represent the current state of the environment as a 36 dimensional vector of rank 1. Here each component of the vector has one of 3 values which indicate:

- 0.5: Agent location
- 1.0: Negative reward
- 0.99: Positive reward
- 1.0: Goal square

We feed this vector to a Feed-forward Network that will give us the estimated action-values. Then We take the argmax over these action-values. We consider the environments as solved when we have achieved a 100 episode running average of 2.1.

8.2 CartPole-v1 & LunarLander-v2

For both these environments we have rather similar action spaces (Discrete), hence the Feed-forward Network architecture used is also similar. The input of the network has the shape $1 \times$ state-space output of network will $1 \times$ action-space.

The environments are considered solved when:

CartPole-v1 Running average reward of 500 (100 episodes)

LunarLander-v2 Running average reward > 200 (100 episodes)

8.3 Training Discussion




Figure 14: Grid Training

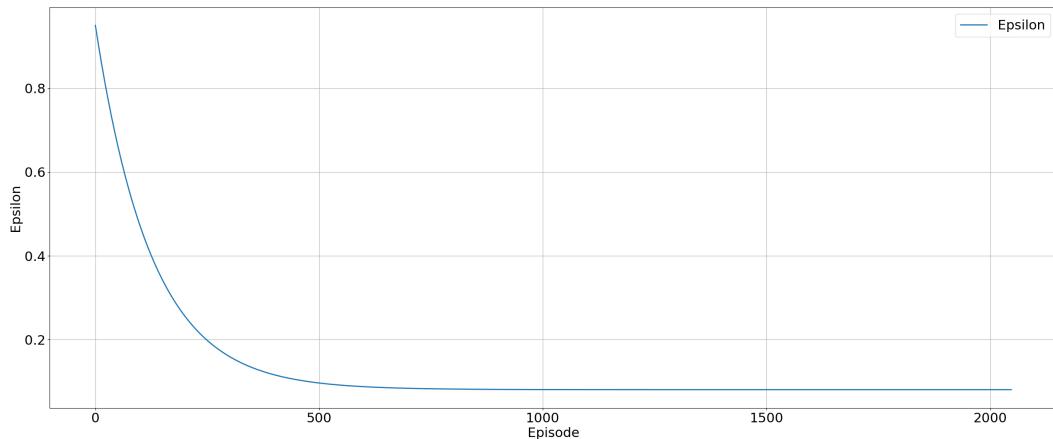


Figure 15: Epsilon Decay – Grid Env

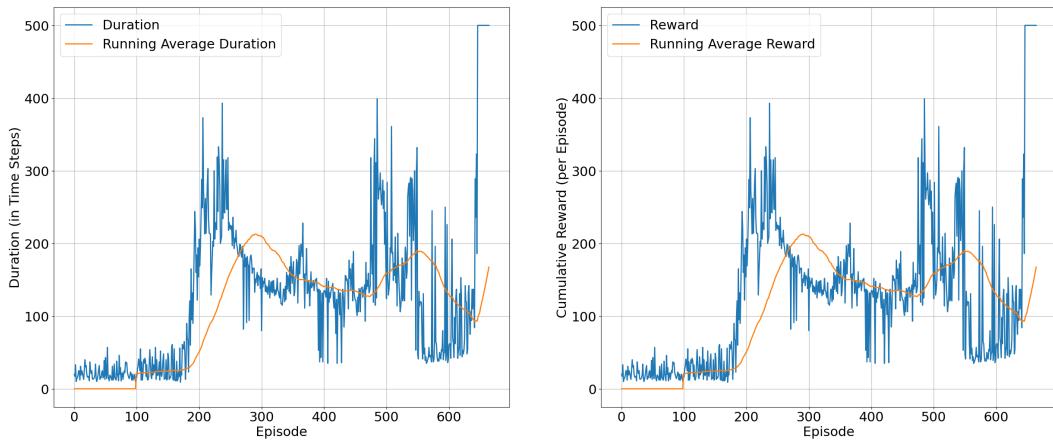


Figure 16: Cart-Pole Training

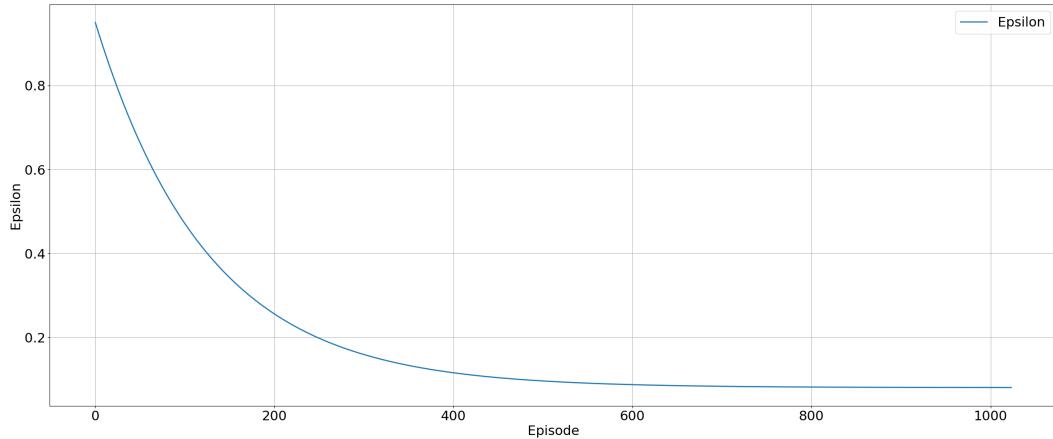


Figure 17: Epsilon Decay – Cart Pole

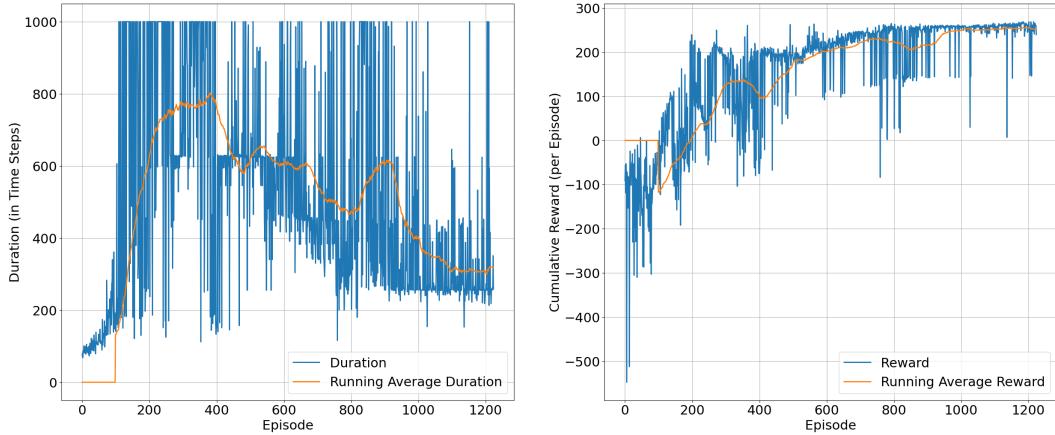


Figure 18: Lunar Lander Training




Figure 19: Epsilon Decay – Lunar Lander

We successfully solved all three environments using simple FeedForward neural networks. Additionally, we implemented a tau-based Target Network update strategy to stabilize the training process and used an experience replay memory of size 10000. Of the three environments, CartPole-v1 exhibited the most unstable learning. This instability is due to the fact that a single missstep in CartPole-v1 can cause the pole to collapse and the episode to end prematurely. The other two environments had more stable learning curves and were successfully solved by the DQN agent.

- 9 Provide the evaluation results. Run your agent on the three environments for at least 5 episodes, where the agent chooses only greedy actions from the learnt policy.**

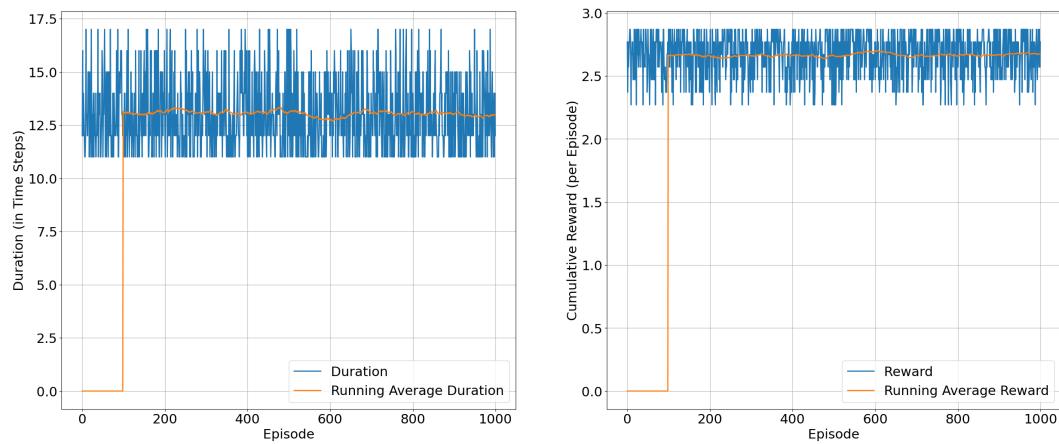


Figure 20: Grid Environment Evaluation




Figure 21: Cart Environment Evaluation

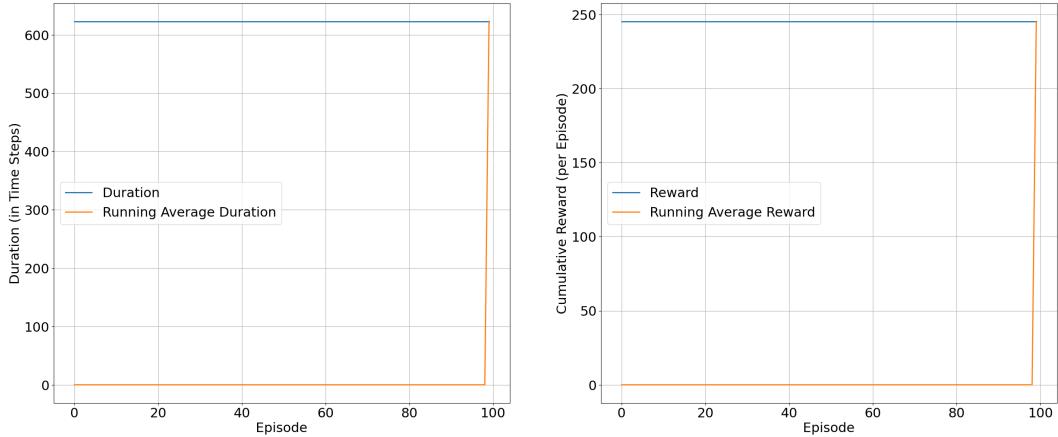


Figure 22: Lunar Environment Evaluation

The Double Q-Learning with a FeedForward neural network approach has successfully learned the optimal policy for all three environments. The evaluation result for the GridEnvironment varies due to the agent being placed at a random start squares, which affect the number of steps required to reach the goal. To encourage the agent to reach the goal as quickly as possible, we applied a negative reward of -0.1 for each step taken.

10 Training Strategy Comparison

10.1 GridWorld

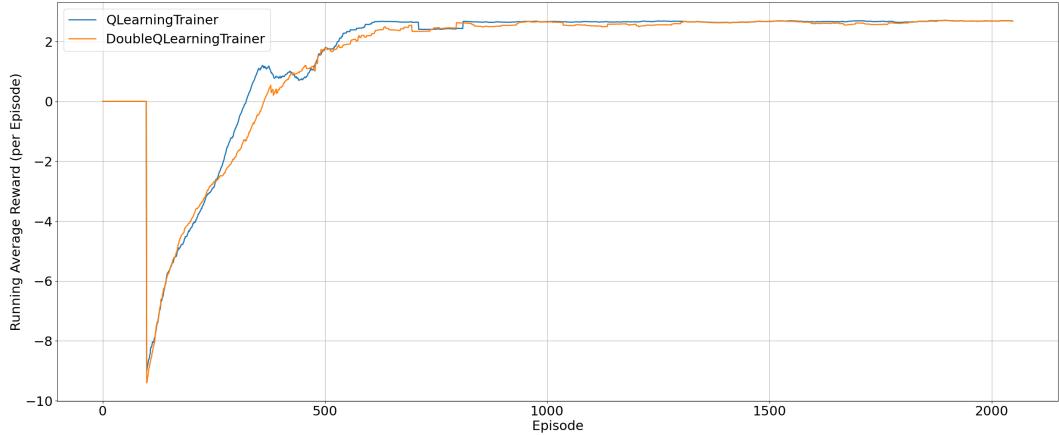


Figure 23: Gridworld Agent Training Comparison

The GridWorld environment is relatively simple, both the DQN and DDQN models may have enough capacity to learn the optimal solution. The capacity of a model refers to its ability to capture the complexities and patterns in the data. In more complex environments, the DDQN model may have an advantage over the DQN model since it can more effectively learn the optimal action values by reducing overestimation bias.

Secondly, since the GridWorld environment is not very complex, the exploration-exploitation trade-off is not a significant concern. Both the DQN and DDQN models can effectively explore the environment and learn the optimal solution without getting stuck in local minima.

Therefore, it is not surprising that the DQN and DDQN models converge at nearly similar rates in a simplistic environment like the GridWorld environment.

10.2 Cartpole v-1




Figure 24: LunarLander Agent Training Comparison

In some instances of this environment, the DQN model may outperform the DDQN model.

One reason for this could be that the DDQN model is highly sensitive to its hyperparameters, such as the learning rate and discount factor. Finding optimal hyperparameters for the DDQN model can be difficult, especially in more complex environments. In contrast, the DQN model may be more stable during training and may require less tuning of hyperparameters.

Overall, the DQN model may be a simpler and more stable choice for solving the CartPole environment, although the performance of both models can depend on the specific instance of the environment and the hyperparameters used.

10.3 LunarLander v-2




Figure 25: LunarLander Agent Training Comparison

The LunarLander v-2 is again rather simple in terms of its state and action spaces. And we see that the DQN slightly outperforms the DDQN algorithm. However, hyperparameters could be a

factor—the performance of both DQN and DDQN can be highly sensitive to their hyperparameters, such as learning rate, batch size, and discount factor. It’s possible that the hyperparameters used for the DQN algorithm were more suitable for the specific instance of the LunarLander v2 environment (discrete), while the hyperparameters used for the DDQN algorithm were not as optimal.

References

- [1] Andrew G Barto, Richard S Sutton, and Charles W Anderson. “Neuronlike adaptive elements that can solve difficult learning control problems”. In: *IEEE transactions on systems, man, and cybernetics* 5 (1983), pp. 834–846.
- [2] Greg Brockman et al. *OpenAI Gym*. 2016. eprint: arXiv:1606.01540.
- [3] Volodymyr Mnih et al. “Playing atari with deep reinforcement learning”. In: *arXiv preprint arXiv:1312.5602* (2013).
- [4] Hado Van Hasselt, Arthur Guez, and David Silver. “Deep reinforcement learning with double q-learning”. In: *Proceedings of the AAAI conference on artificial intelligence*. Vol. 30. 1. 2016.